

# Homework 3

Kris Hanus, Laura Glathar, Arkya Rakshit, Jace Crist, Brandon Dlugosz  
University of Nebraska at Omaha

BSAD 8700 - Business Analytics  
Due: February 2, 2015

## ANSWER FOR 10:

```
(a) tail(Weekly, 1)

##      Year   Lag1   Lag2   Lag3   Lag4   Lag5   Volume Today Direction
## 1089 2010 1.034 0.283 1.281 2.969 -0.861 2.707105 0.069       Up

summary(Weekly)

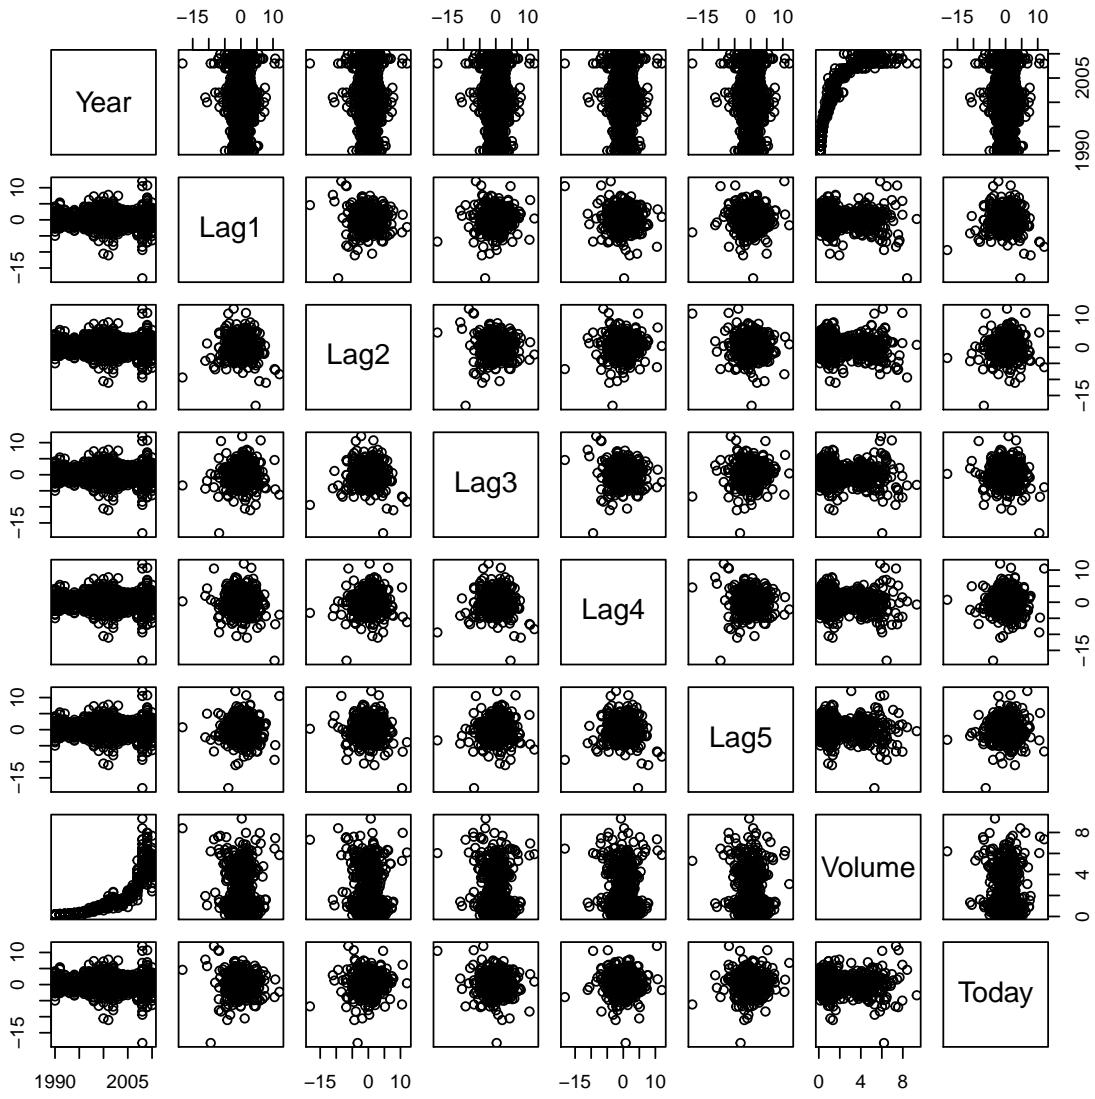
##      Year           Lag1           Lag2           Lag3
## Min. :1990  Min. :-18.1950  Min. :-18.1950  Min. :-18.1950
## 1st Qu.:1995 1st Qu.:-1.1540  1st Qu.:-1.1540  1st Qu.:-1.1580
## Median :2000 Median : 0.2410  Median : 0.2410  Median : 0.2410
## Mean   :2000  Mean  : 0.1506  Mean  : 0.1511  Mean  : 0.1472
## 3rd Qu.:2005 3rd Qu.: 1.4050  3rd Qu.: 1.4090  3rd Qu.: 1.4090
## Max.   :2010  Max.  :12.0260  Max.  :12.0260  Max.  :12.0260
##           Lag4           Lag5           Volume
## Min. :-18.1950  Min. :-18.1950  Min. :0.08747
## 1st Qu.:-1.1580  1st Qu.:-1.1660  1st Qu.:0.33202
## Median : 0.2380  Median : 0.2340  Median :1.00268
## Mean   : 0.1458  Mean  : 0.1399  Mean  :1.57462
## 3rd Qu.: 1.4090  3rd Qu.: 1.4050  3rd Qu.:2.05373
## Max.   :12.0260  Max.  :12.0260  Max.  :9.32821
##           Today          Direction
## Min. :-18.1950  Down:484
## 1st Qu.:-1.1540  Up  :605
## Median : 0.2410
## Mean   : 0.1499
## 3rd Qu.: 1.4050
## Max.   :12.0260

data1<-Weekly[,1:8]
attach(Weekly)
cor(data1)

##             Year           Lag1           Lag2           Lag3           Lag4
## Year 1.000000000 -0.032289274 -0.03339001 -0.03000649 -0.031127923
## Lag1 -0.032289274 1.000000000 -0.07485305  0.05863568 -0.071273876
## Lag2 -0.03339001 -0.074853051 1.000000000 -0.07572091  0.058381535
```

```
## Lag3 -0.03000649  0.058635682 -0.07572091  1.000000000 -0.075395865
## Lag4 -0.03112792 -0.071273876  0.05838153 -0.07539587  1.000000000
## Lag5 -0.03051910 -0.008183096 -0.07249948  0.06065717 -0.075675027
## Volume 0.84194162 -0.064951313 -0.08551314 -0.06928771 -0.061074617
## Today -0.03245989 -0.075031842  0.05916672 -0.07124364 -0.007825873
##          Lag5      Volume     Today
## Year    -0.030519101 0.84194162 -0.032459894
## Lag1    -0.008183096 -0.06495131 -0.075031842
## Lag2    -0.072499482 -0.08551314  0.059166717
## Lag3     0.060657175 -0.06928771 -0.071243639
## Lag4    -0.075675027 -0.06107462 -0.007825873
## Lag5     1.000000000 -0.05851741  0.011012698
## Volume -0.058517414  1.000000000 -0.033077783
## Today   0.011012698 -0.03307778  1.000000000

pairs(data1)
```



There are a few interesting places of correlation. Primarily, with Volume and Year. Other wise, it is observable that each of the Lags are clustered, but it is difficult to observe other relationships.

```
(b) glm.fit=glm(Direction~Lag1+Lag2+Lag3+Lag4+Lag5+Volume,data=Weekly, family = binomial)
summary(glm.fit)

##
## Call:
## glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
##       Volume, family = binomial, data = Weekly)
##
## Deviance Residuals:
##      Min        1Q    Median        3Q       Max
## -1.6949   -1.2565    0.9913    1.0849    1.4579
##
```

```

## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) 0.26686   0.08593   3.106   0.0019 **
## Lag1        -0.04127   0.02641  -1.563   0.1181
## Lag2         0.05844   0.02686   2.175   0.0296 *
## Lag3        -0.01606   0.02666  -0.602   0.5469
## Lag4        -0.02779   0.02646  -1.050   0.2937
## Lag5        -0.01447   0.02638  -0.549   0.5833
## Volume      -0.02274   0.03690  -0.616   0.5377
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 1496.2 on 1088 degrees of freedom
## Residual deviance: 1486.4 on 1082 degrees of freedom
## AIC: 1500.4
##
## Number of Fisher Scoring iterations: 4

coef(glm.fit)

## (Intercept)      Lag1      Lag2      Lag3      Lag4      Lag5
## 0.26686414 -0.04126894  0.05844168 -0.01606114 -0.02779021 -0.01447206
##       Volume
## -0.02274153

```

The only predictors which have significance are the intercept and Lag2. Lag2 is between 95% and 99% significant. The Intercept is 99% and 99.9%.

```

(c) contrasts(Direction)

##      Up
## Down  0
## Up    1

glm.pred=rep("Down", 1089)
glm.probs=predict(glm.fit,type="response")
glm.probs[1:10]

##      1      2      3      4      5      6      7
## 0.6086249 0.6010314 0.5875699 0.4816416 0.6169013 0.5684190 0.5786097
##      8      9     10
## 0.5151972 0.5715200 0.5554287

glm.pred[glm.probs>0.5] <- "Up"
table(glm.pred,Direction)

##          Direction
## glm.pred Down Up
##      Down    54 48
##      Up     430 557

```

```
557/(557+430)
```

```
## [1] 0.5643364
```

```
430/(557+430)
```

```
## [1] 0.4356636
```

```
48/(48+54)
```

```
## [1] 0.4705882
```

The confusion matrix shows that on days when logistic regression predicts an increase in the market, it has a 56.4% accuracy rate. The error rate is 43.6% for predicting Up and is actually Down. The error rate is 47.1% for predicting Down and is actually Up.

```
(d) glm.fit2=glm(Direction~Lag2, data=Weekly, family = binomial)
summary(glm.fit2)

##
## Call:
## glm(formula = Direction ~ Lag2, family = binomial, data = Weekly)
##
## Deviance Residuals:
##      Min     1Q   Median     3Q    Max 
## -1.564 -1.267  1.008  1.086  1.386 
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)    
## (Intercept) 0.21473   0.06123   3.507 0.000453 ***
## Lag2        0.06279   0.02636   2.382 0.017230 *  
## ---        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 1496.2 on 1088 degrees of freedom
## Residual deviance: 1490.4 on 1087 degrees of freedom
## AIC: 1494.4
##
## Number of Fisher Scoring iterations: 4

coef(glm.fit2)

## (Intercept)      Lag2
## 0.21473151  0.06279058

contrasts(Direction)

##      Up
## Down  0
## Up    1
```

```

glm.pred2=rep("Down", 1089)
glm.probs2=predict(glm.fit2,type="response")
glm.probs2[1:10]

##      1       2       3       4       5       6       7
## 0.5777243 0.5661029 0.5492840 0.5132426 0.6071571 0.5644982 0.5716776
##      8       9      10
## 0.5321015 0.5659641 0.5541137

glm.pred2[glm.probs2>0.5] <- "Up"
table(glm.pred2,Direction)

##          Direction
##  glm.pred2 Down Up
##      Down    33   26
##      Up      451  579

579/(579+451)

## [1] 0.5621359

451/(579+451)

## [1] 0.4378641

26/(26+33)

## [1] 0.440678

```

The confusion matrix shows that on days when logistic regression predicts an increase in the market, it has a 56.2% accuracy rate. The error rate is 43.8% for predicting Up and is actually Down. The error rate is 44.1% for predicting Down and is actually Up.

#### ANSWER FOR 11:

```

(a) #install.packages("RCurl")
library(RCurl)

## Loading required package: bitops

dataset<-getURL(
  'https://raw.githubusercontent.com/Jwcrist/BusA/master/homeworks/Assignment%204/Auto.csv',
  ssl.verifypeer=0L, followlocation=1L)
dataset1<-read.csv(text=dataset)
head(dataset1,3)

##   mpg cylinders displacement horsepower weight acceleration year origin
## 1 18         8           307        130   3504         12.0     70      1
## 2 15         8           350        165   3693         11.5     70      1
## 3 18         8           318        150   3436         11.0     70      1
##                               name mpg01

```

```

## 1 chevrolet chevelle malibu      0
## 2          buick skylark 320      0
## 3        plymouth satellite      0

```

As can be shown above we have created our binary data and a median column

```

(b) library(dplyr)
## Warning: package 'dplyr' was built under R version 3.1.2
##
## Attaching package: 'dplyr'
##
## The following object is masked from 'package:MASS':
##
##     select
##
## The following object is masked from 'package:stats':
##
##     filter
##
## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union

dataset2<-dataset1 %>%
  select(mpg, cylinders, displacement, as.numeric(horsepower), weight, acceleration, year, origin, m
str(dataset2)

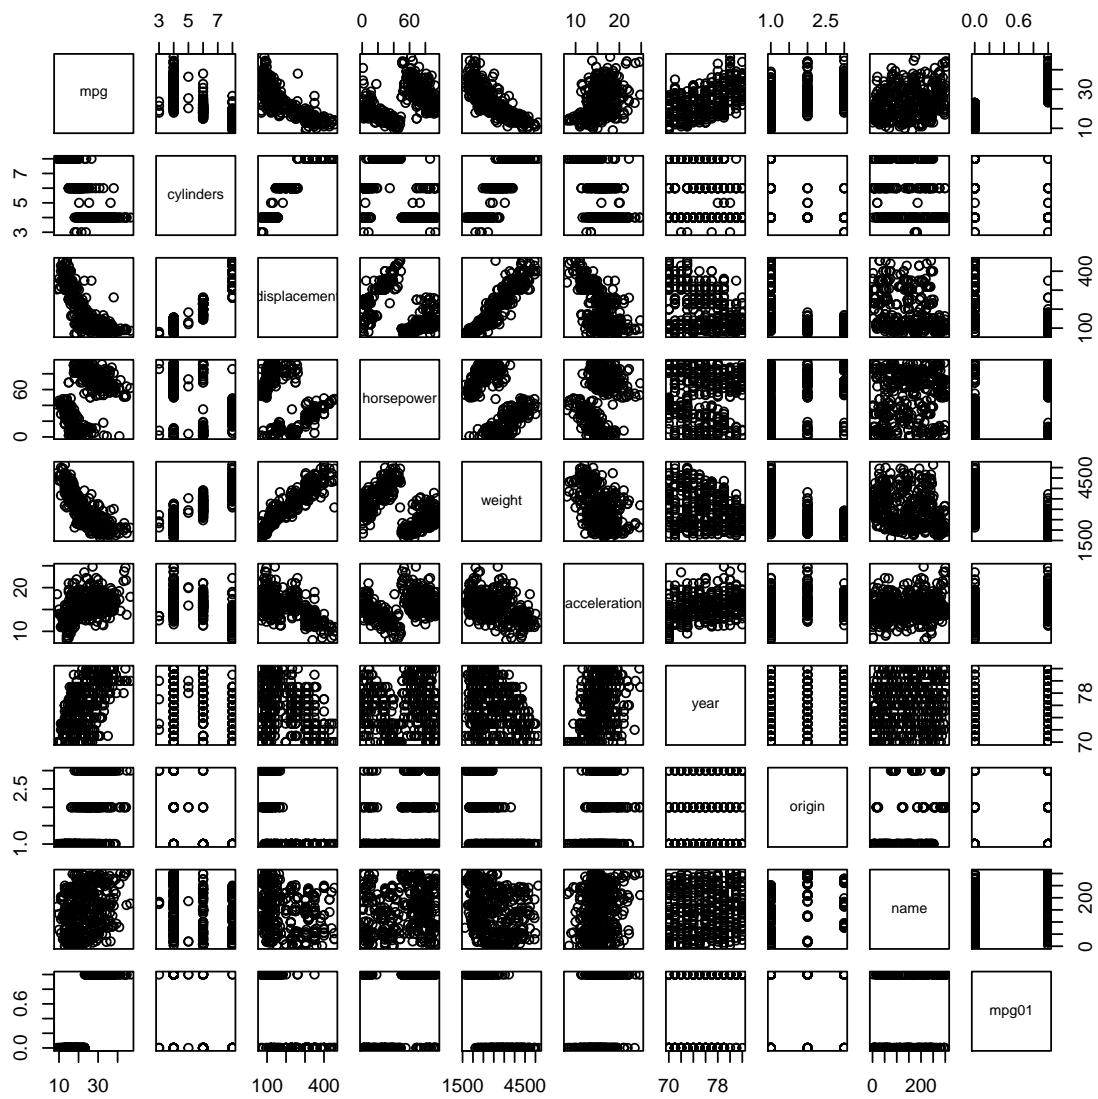
## 'data.frame': 397 obs. of  9 variables:
## $ mpg           : num  18 15 18 16 17 15 14 14 14 15 ...
## $ cylinders    : int  8 8 8 8 8 8 8 8 8 8 ...
## $ displacement: num  307 350 318 304 302 429 454 440 455 390 ...
## $ horsepower   : Factor w/ 94 levels "?","100","102",...: 17 35 29 29 24 42 47 46 48 40 ...
## $ weight        : int  3504 3693 3436 3433 3449 4341 4354 4312 4425 3850 ...
## $ acceleration: num  12 11.5 11 12 10.5 10 9 8.5 10 8.5 ...
## $ year         : int  70 70 70 70 70 70 70 70 70 70 ...
## $ origin        : int  1 1 1 1 1 1 1 1 1 1 ...
## $ mpg01        : int  0 0 0 0 0 0 0 0 0 0 ...

cor(dataset2)

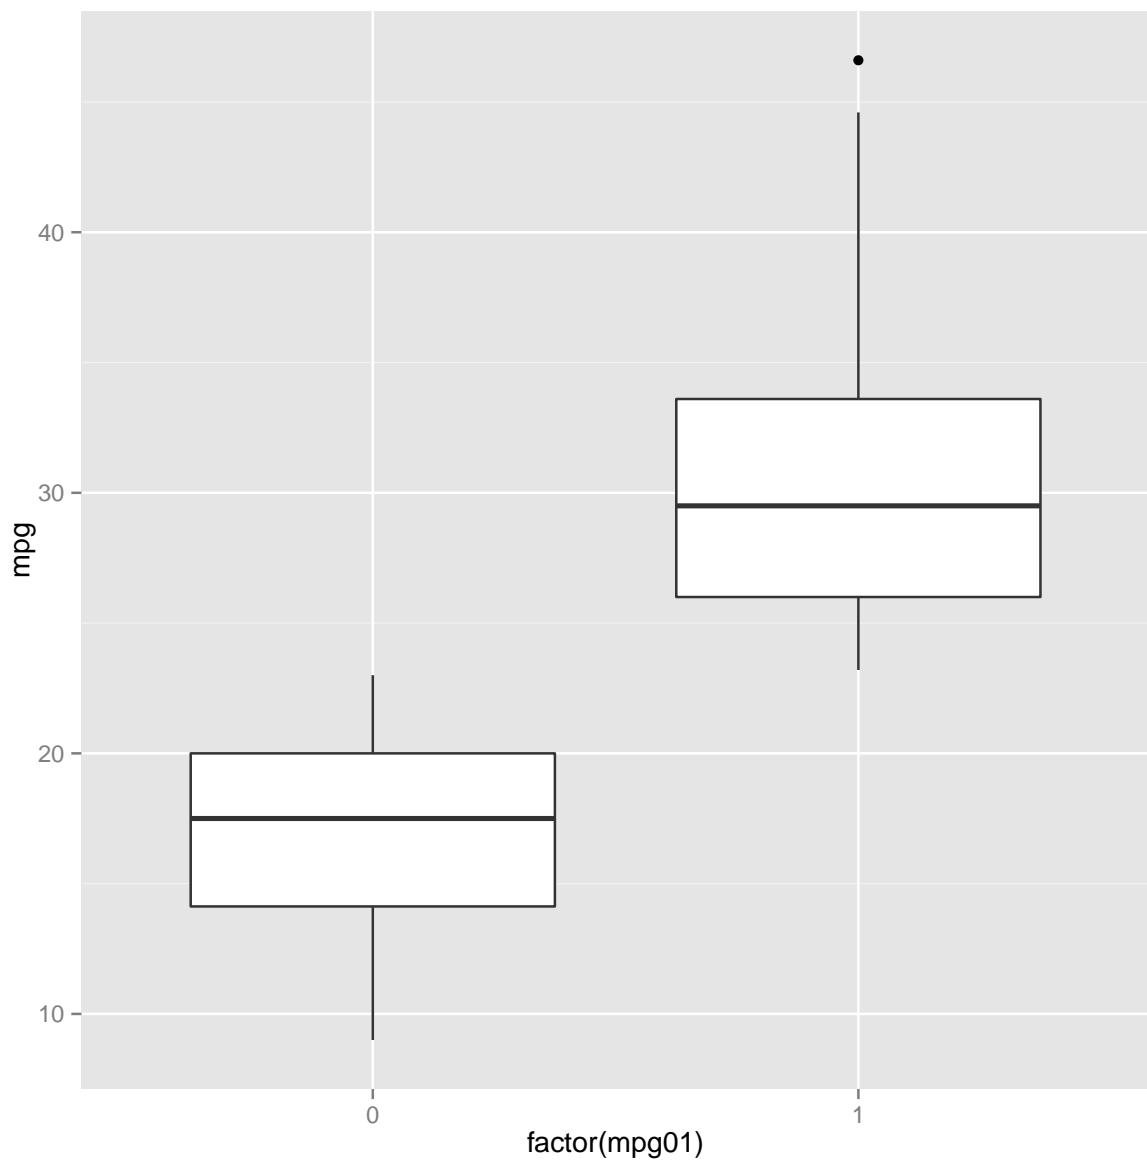
## Error in cor(dataset2): 'x' must be numeric

pairs(dataset1)

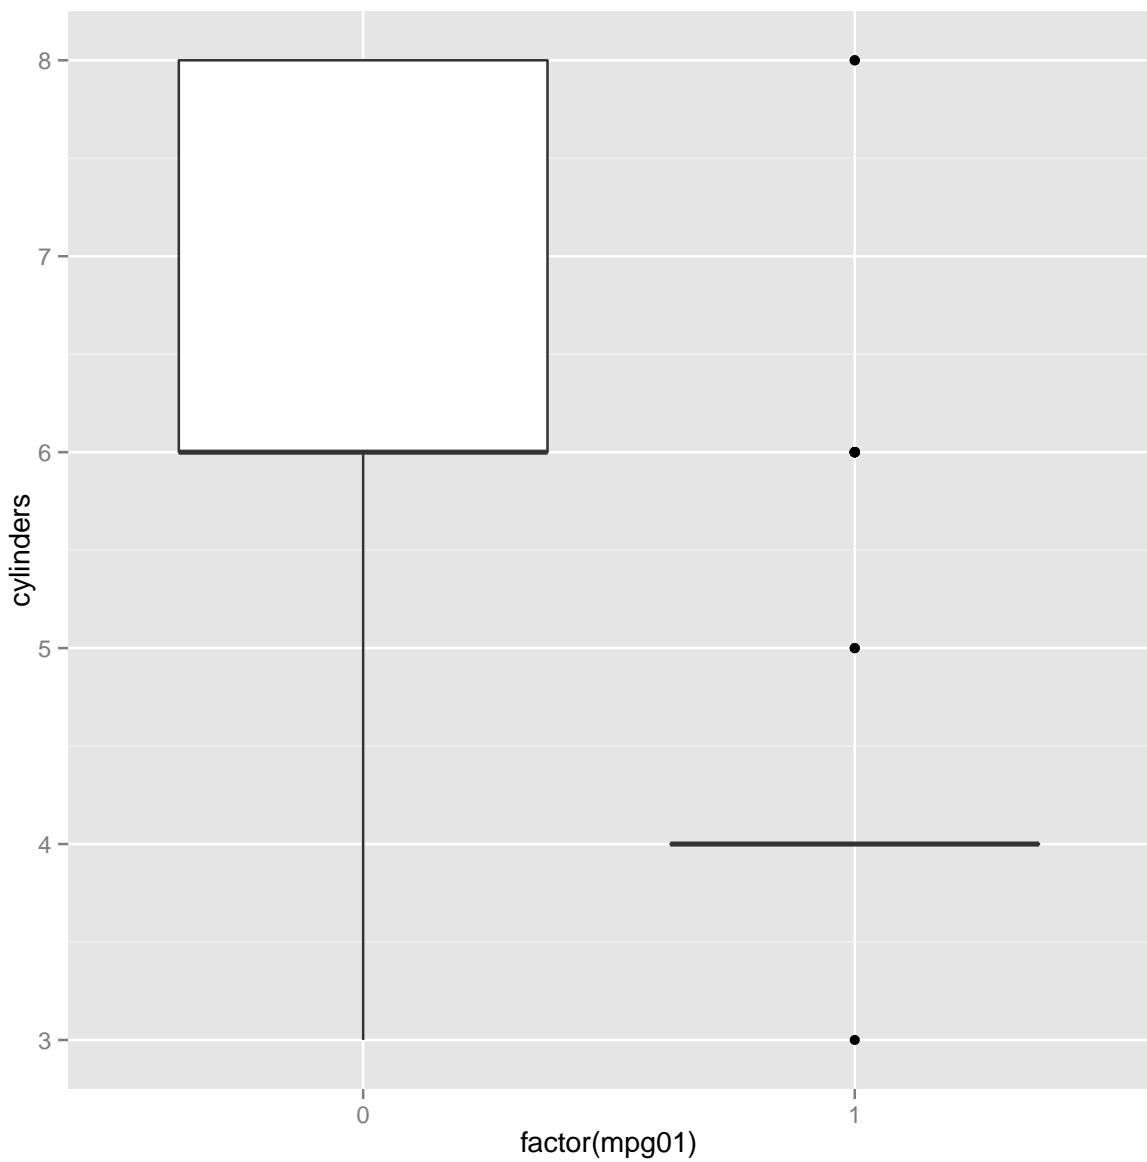
```



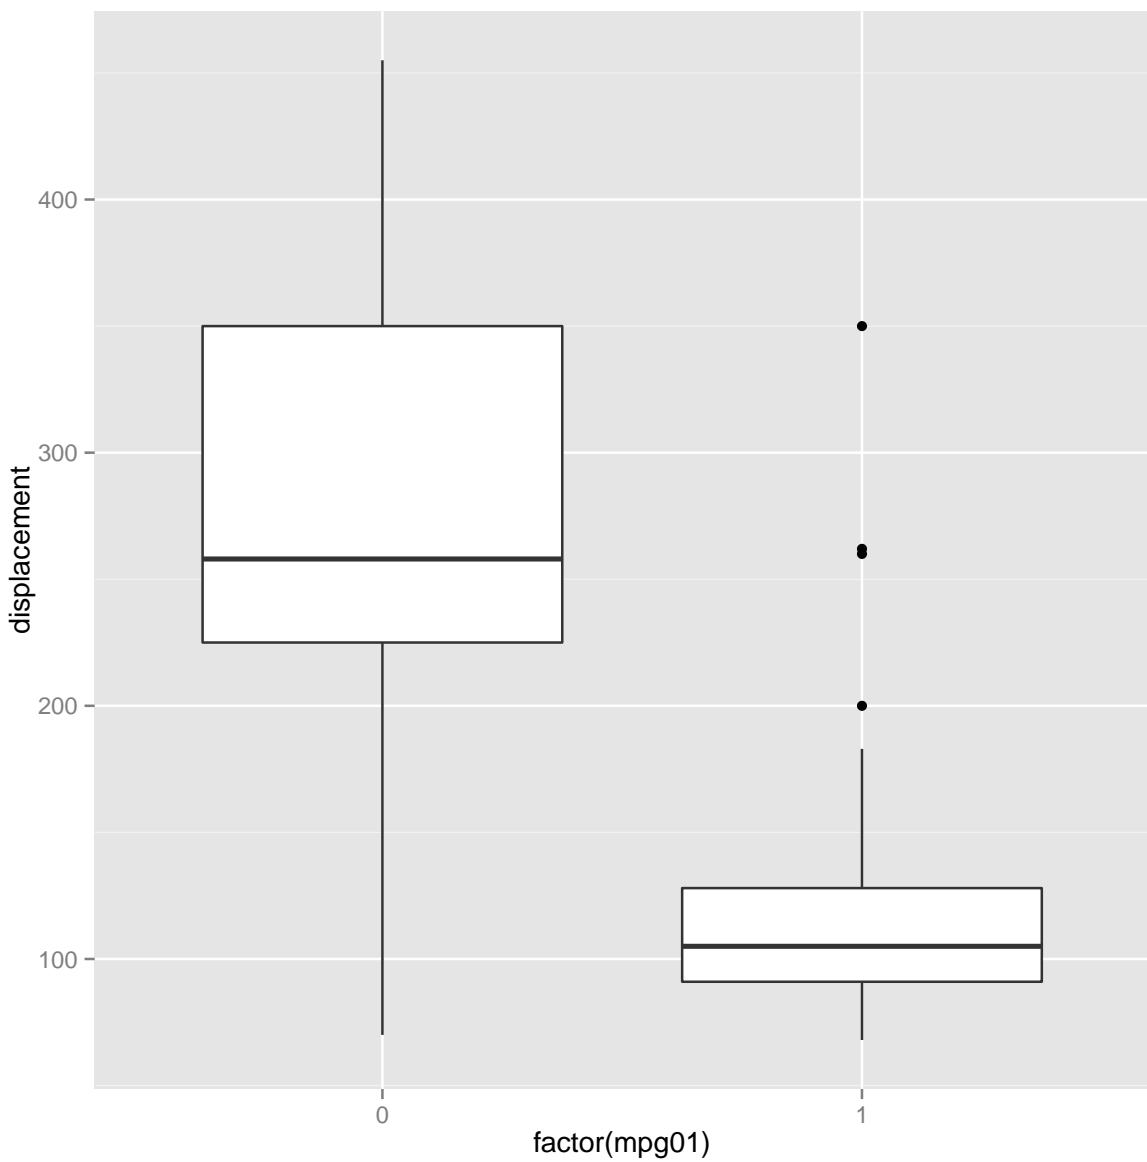
```
library(ggplot2)
ggplot(dataset1, aes(x=factor(mpg01), y=mpg))+geom_boxplot()
```



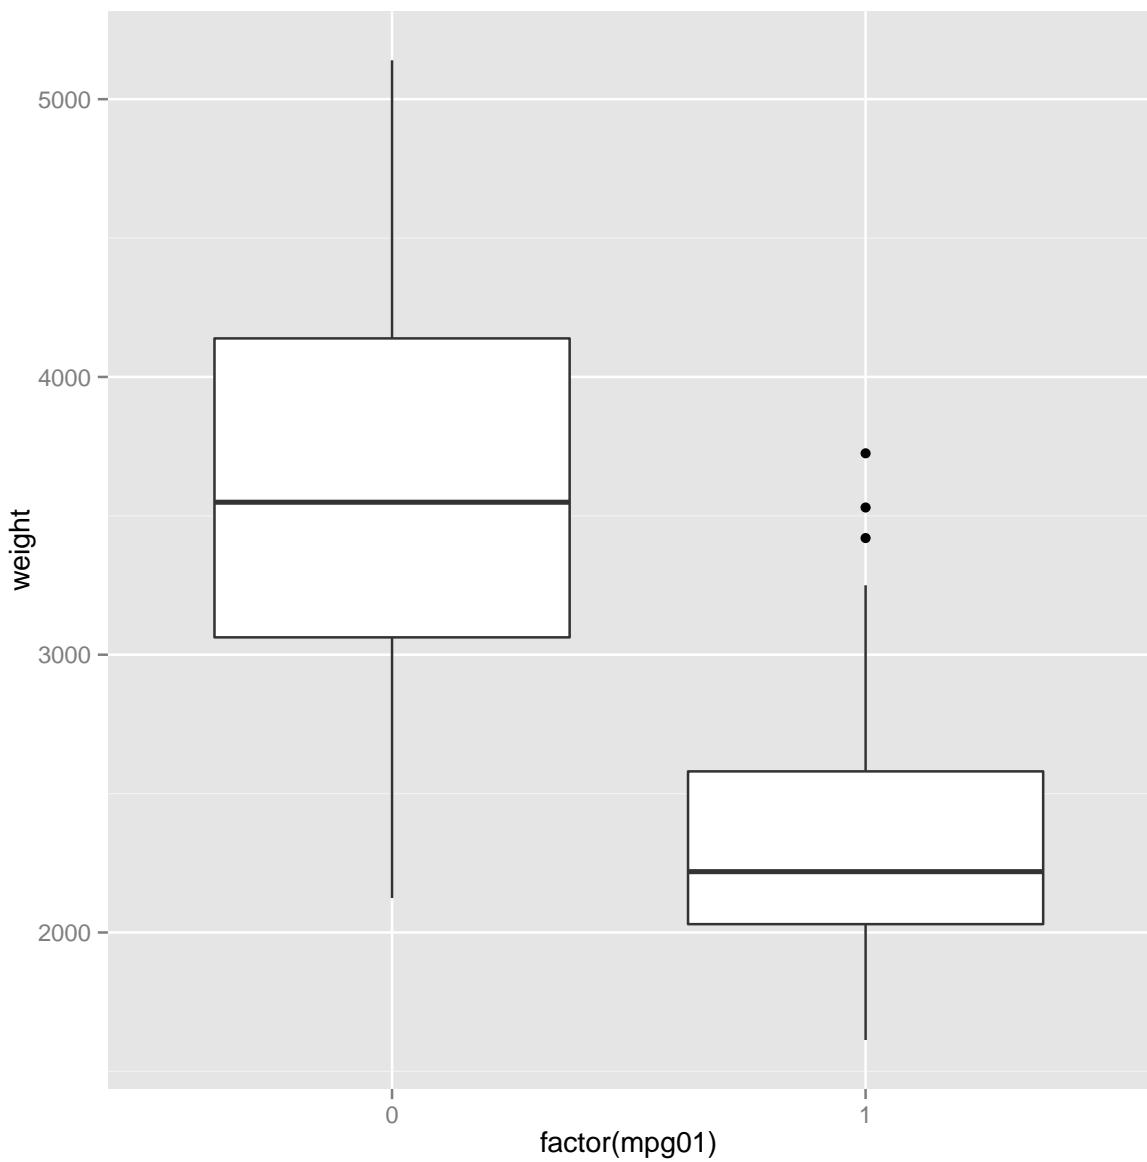
```
ggplot(dataset1, aes(x=factor(mpg01), y=cylinders)) + geom_boxplot()
```



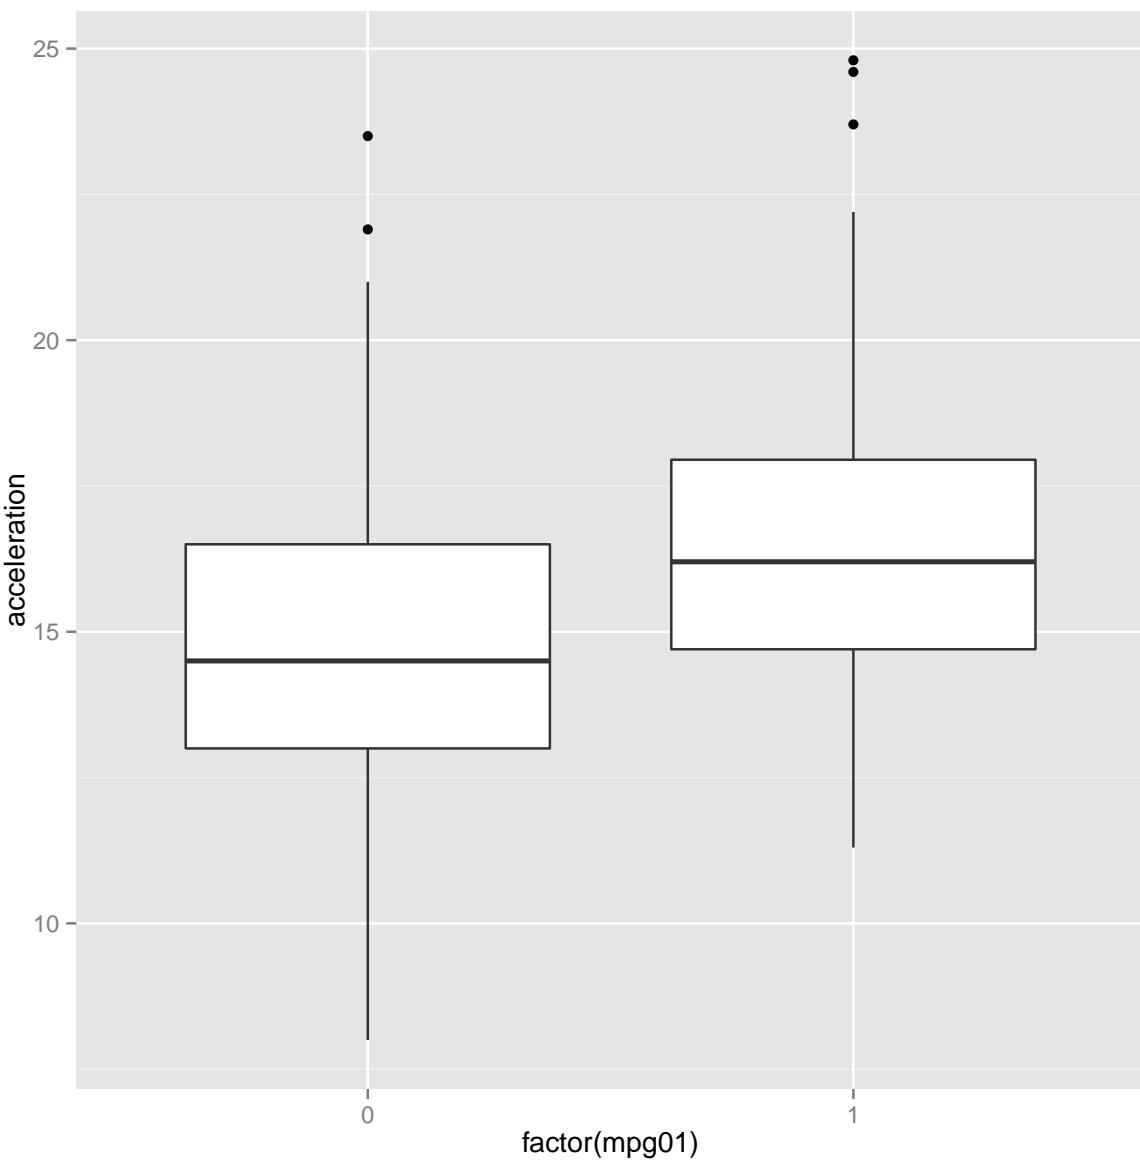
```
ggplot(dataset1, aes(x=factor(mpg01), y=displacement))+geom_boxplot()
```



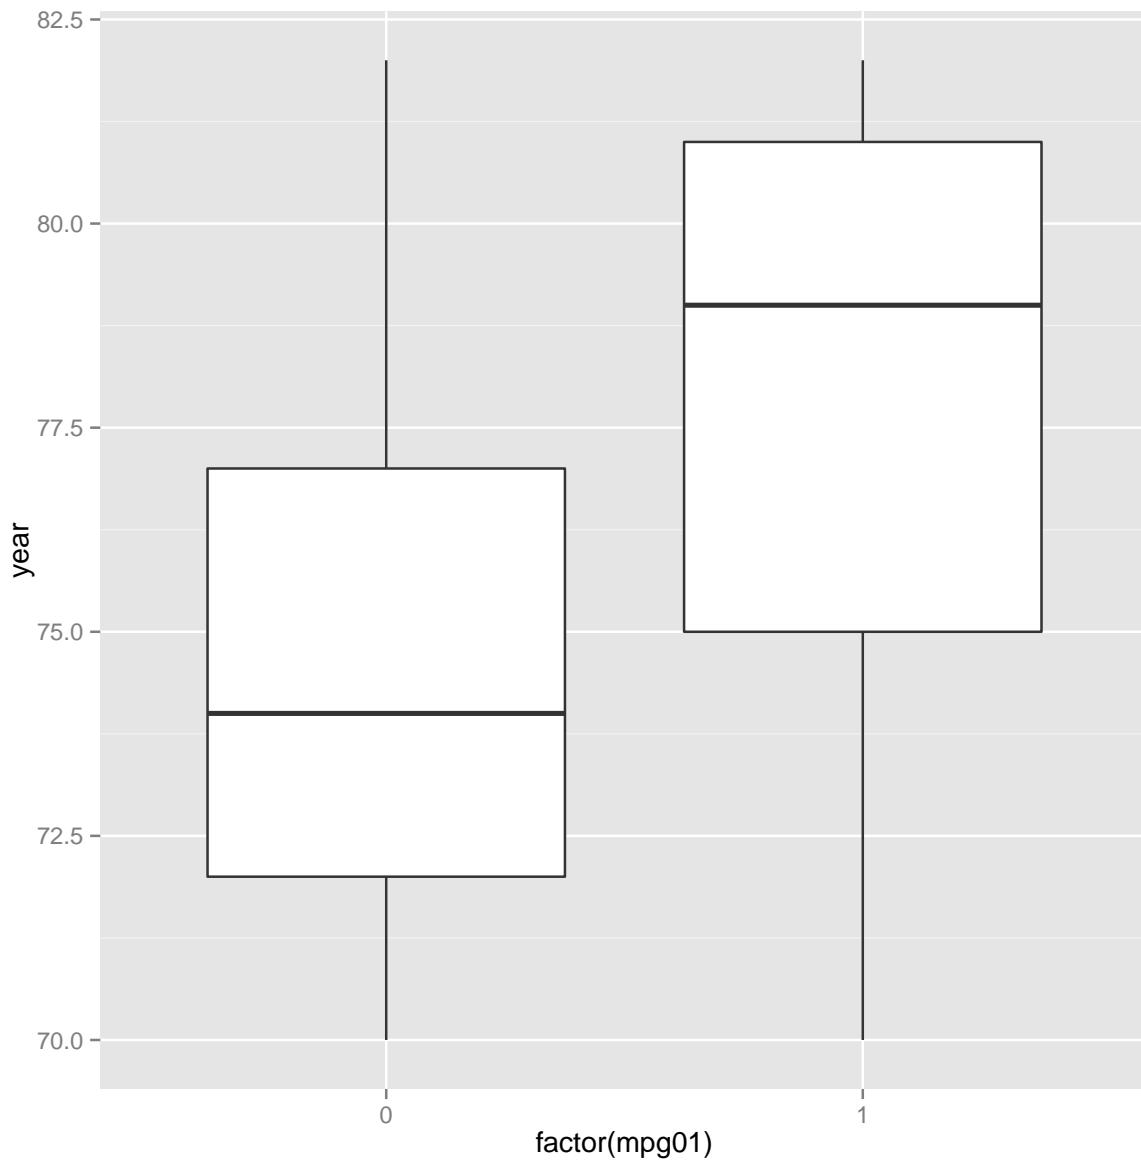
```
ggplot(dataset1, aes(x=factor(mpg01), y=weight)) + geom_boxplot()
```



```
ggplot(dataset1, aes(x=factor(mpg01), y=acceleration))+geom_boxplot()
```



```
ggplot(dataset1, aes(x=factor(mpg01), y=year))+geom_boxplot()
```



There is too many layers to observe anything of significance in horsepower. Acceleration and year may have be significantly different from eachother. All the other plots show that there is probable differences between our binary factors.

(c) `tail(dataset1)`

```
##      mpg cylinders displacement horsepower weight acceleration year origin
## 392    27          4         151        90   2950       17.3     82     1
## 393    27          4         140        86   2790       15.6     82     1
## 394    44          4         97        52   2130       24.6     82     2
## 395    32          4        135        84   2295       11.6     82     1
## 396    28          4        120        79   2625       18.6     82     1
## 397    31          4        119        82   2720       19.4     82     1
##                               name mpg01
## 392  chevrolet camaro     1
```

```

## 393  ford mustang gl      1
## 394      vw pickup      1
## 395      dodge rampage    1
## 396      ford ranger      1
## 397      chevy s-10       1

trainingdata<-dataset1[1:318,]
testdata<-dataset1[79,]

```

A 80% training data set was chosen to be used in order to prove against our test data.

(d) Do Not Do

(e) Do Not Do

```

(f) glm.fit3=glm(mpg01~Lag2,data=Weekly, family = binomial)
## Error in eval(expr, envir, enclos): object 'mpg01' not found

summary(glm.fit2)

##
## Call:
## glm(formula = Direction ~ Lag2, family = binomial, data = Weekly)
##
## Deviance Residuals:
##     Min      1Q  Median      3Q      Max
## -1.564  -1.267   1.008   1.086   1.386
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.21473   0.06123   3.507 0.000453 ***
## Lag2         0.06279   0.02636   2.382 0.017230 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 1496.2 on 1088 degrees of freedom
## Residual deviance: 1490.4 on 1087 degrees of freedom
## AIC: 1494.4
##
## Number of Fisher Scoring iterations: 4

coef(glm.fit2)

## (Intercept)      Lag2
## 0.21473151  0.06279058

contrasts(Direction)

##      Up
## Down  0
## Up    1

```

```

glm.pred2=rep("Down", 1089)
glm.probs2=predict(glm.fit2,type="response")
glm.probs2[1:10]

##      1       2       3       4       5       6       7
## 0.5777243 0.5661029 0.5492840 0.5132426 0.6071571 0.5644982 0.5716776
##      8       9      10
## 0.5321015 0.5659641 0.5541137

glm.pred2[glm.probs2>0.5] <- "Up"
table(glm.pred2,Direction)

##          Direction
##  glm.pred2 Down  Up
##      Down    33   26
##      Up     451  579

```