

Homework 3

Kris Hanus, Laura Glathar, Arkya Rakshit, Jace Crist, Brandon Dlugosz
University of Nebraska at Omaha

BSAD 8700 - Business Analytics
Due: February 2, 2015

ANSWER FOR 10:

```
(a) tail(Weekly, 1)

##      Year   Lag1   Lag2   Lag3   Lag4   Lag5   Volume Today Direction
## 1089 2010 1.034 0.283 1.281 2.969 -0.861 2.707105 0.069       Up

summary(Weekly)

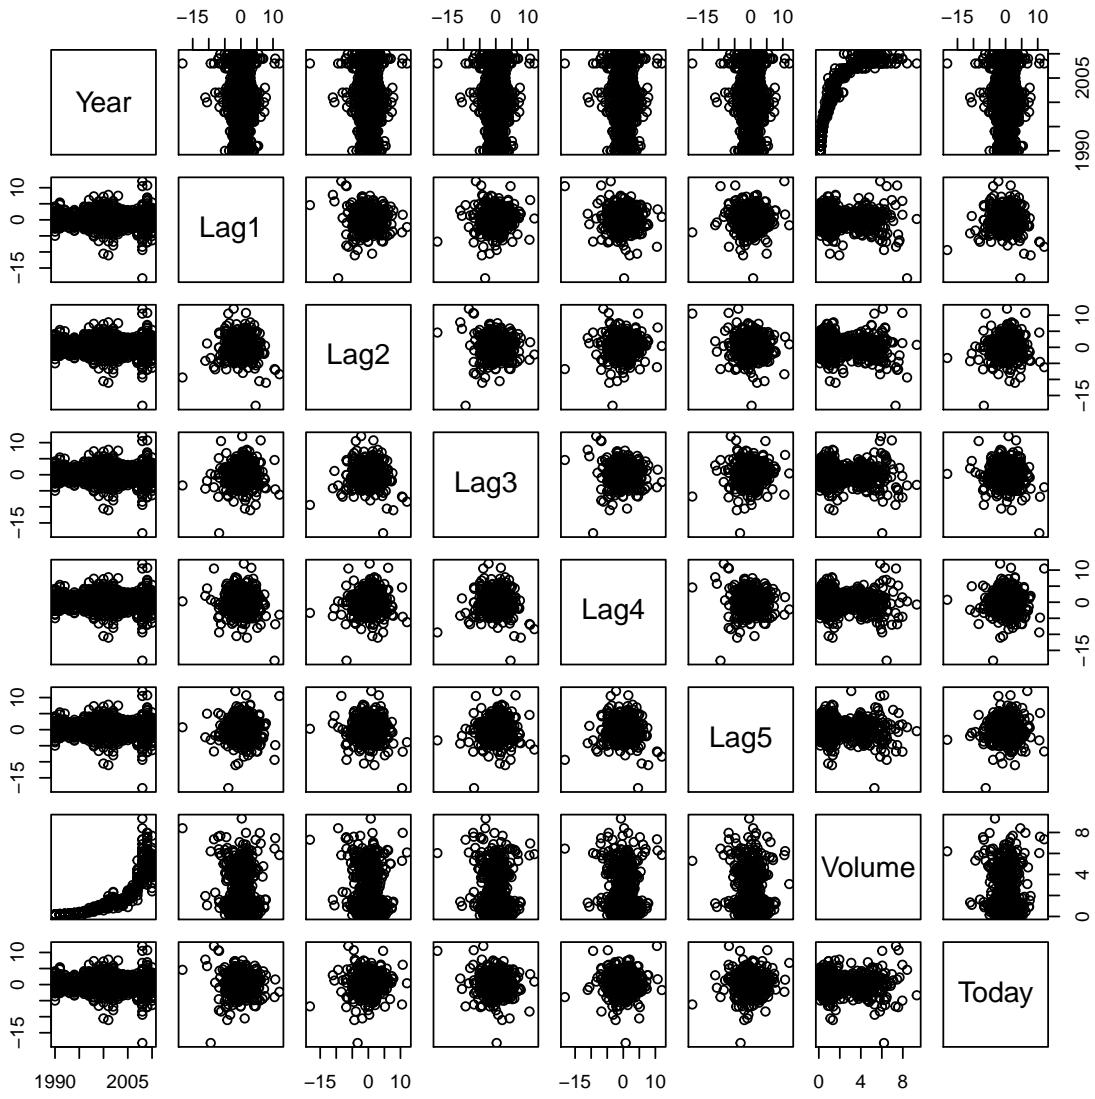
##      Year          Lag1          Lag2          Lag3
##  Min.   :1990   Min.   :-18.1950   Min.   :-18.1950   Min.   :-18.1950
##  1st Qu.:1995   1st Qu.: -1.1540   1st Qu.: -1.1540   1st Qu.: -1.1580
##  Median :2000   Median :  0.2410   Median :  0.2410   Median :  0.2410
##  Mean   :2000   Mean   :  0.1506   Mean   :  0.1511   Mean   :  0.1472
##  3rd Qu.:2005   3rd Qu.:  1.4050   3rd Qu.:  1.4090   3rd Qu.:  1.4090
##  Max.   :2010   Max.   : 12.0260   Max.   : 12.0260   Max.   : 12.0260
##          Lag4          Lag5          Volume
##  Min.   :-18.1950   Min.   :-18.1950   Min.   :0.08747
##  1st Qu.: -1.1580   1st Qu.: -1.1660   1st Qu.:0.33202
##  Median :  0.2380   Median :  0.2340   Median :1.00268
##  Mean   :  0.1458   Mean   :  0.1399   Mean   :1.57462
##  3rd Qu.:  1.4090   3rd Qu.:  1.4050   3rd Qu.:2.05373
##  Max.   : 12.0260   Max.   : 12.0260   Max.   :9.32821
##          Today          Direction
##  Min.   :-18.1950   Down:484
##  1st Qu.: -1.1540   Up  :605
##  Median :  0.2410
##  Mean   :  0.1499
##  3rd Qu.:  1.4050
##  Max.   : 12.0260

data1<-Weekly[,1:8]
attach(Weekly)
cor(data1)

##           Year          Lag1          Lag2          Lag3          Lag4
##  Year  1.00000000 -0.032289274 -0.03339001 -0.03000649 -0.031127923
##  Lag1 -0.03228927  1.000000000 -0.07485305  0.05863568 -0.071273876
##  Lag2 -0.03339001 -0.074853051  1.000000000 -0.07572091  0.058381535
```

```
## Lag3 -0.03000649  0.058635682 -0.07572091  1.000000000 -0.075395865
## Lag4 -0.03112792 -0.071273876  0.05838153 -0.07539587  1.000000000
## Lag5 -0.03051910 -0.008183096 -0.07249948  0.06065717 -0.075675027
## Volume 0.84194162 -0.064951313 -0.08551314 -0.06928771 -0.061074617
## Today -0.03245989 -0.075031842  0.05916672 -0.07124364 -0.007825873
##          Lag5      Volume     Today
## Year    -0.030519101 0.84194162 -0.032459894
## Lag1    -0.008183096 -0.06495131 -0.075031842
## Lag2    -0.072499482 -0.08551314  0.059166717
## Lag3     0.060657175 -0.06928771 -0.071243639
## Lag4    -0.075675027 -0.06107462 -0.007825873
## Lag5     1.000000000 -0.05851741  0.011012698
## Volume -0.058517414  1.000000000 -0.033077783
## Today   0.011012698 -0.03307778  1.000000000

pairs(data1)
```



There are a few interesting places of correlation. Primarily, with Volume and Year. Other wise, it is observable that each of the Lags are clustered, but it is difficult to observe other relationships.

```
(b) glm.fit=glm(Direction~Lag1+Lag2+Lag3+Lag4+Lag5+Volume,data=Weekly, family = binomial)
summary(glm.fit)

##
## Call:
## glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
##     Volume, family = binomial, data = Weekly)
##
## Deviance Residuals:
##      Min        1Q    Median        3Q       Max
## -1.6949   -1.2565    0.9913    1.0849    1.4579
##
```

```

## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) 0.26686   0.08593   3.106   0.0019 **
## Lag1        -0.04127   0.02641  -1.563   0.1181
## Lag2         0.05844   0.02686   2.175   0.0296 *
## Lag3        -0.01606   0.02666  -0.602   0.5469
## Lag4        -0.02779   0.02646  -1.050   0.2937
## Lag5        -0.01447   0.02638  -0.549   0.5833
## Volume      -0.02274   0.03690  -0.616   0.5377
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 1496.2 on 1088 degrees of freedom
## Residual deviance: 1486.4 on 1082 degrees of freedom
## AIC: 1500.4
##
## Number of Fisher Scoring iterations: 4

coef(glm.fit)

## (Intercept)      Lag1      Lag2      Lag3      Lag4      Lag5
## 0.26686414 -0.04126894  0.05844168 -0.01606114 -0.02779021 -0.01447206
##       Volume
## -0.02274153

```

The only predictors which have significance are the intercept and Lag2. Lag2 is between 95% and 99% significant. The Intercept is 99% and 99.9%.

```

(c) contrasts(Direction)

##      Up
## Down  0
## Up    1

glm.pred=rep("Down", 1089)
glm.probs=predict(glm.fit,type="response")
glm.probs[1:10]

##      1      2      3      4      5      6      7
## 0.6086249 0.6010314 0.5875699 0.4816416 0.6169013 0.5684190 0.5786097
##      8      9     10
## 0.5151972 0.5715200 0.5554287

glm.pred[glm.probs>0.5] <- "Up"
table(glm.pred,Direction)

##          Direction
## glm.pred Down Up
##      Down    54 48
##      Up     430 557

```

```
557/(557+430)
```

```
## [1] 0.5643364
```

```
430/(557+430)
```

```
## [1] 0.4356636
```

```
48/(48+54)
```

```
## [1] 0.4705882
```

The confusion matrix shows that on days when logistic regression predicts an increase in the market, it has a 56.4% accuracy rate. The error rate is 43.6% for predicting Up and is actually Down. The error rate is 47.1% for predicting Down and is actually Up.

```
(d) glm.fit2=glm(Direction~Lag2, data=Weekly, family = binomial)
summary(glm.fit2)

##
## Call:
## glm(formula = Direction ~ Lag2, family = binomial, data = Weekly)
##
## Deviance Residuals:
##      Min     1Q   Median     3Q    Max 
## -1.564 -1.267  1.008  1.086  1.386 
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)    
## (Intercept) 0.21473   0.06123   3.507 0.000453 ***
## Lag2        0.06279   0.02636   2.382 0.017230 *  
## ---        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 1496.2 on 1088 degrees of freedom
## Residual deviance: 1490.4 on 1087 degrees of freedom
## AIC: 1494.4
##
## Number of Fisher Scoring iterations: 4

coef(glm.fit2)

## (Intercept)      Lag2
## 0.21473151  0.06279058

contrasts(Direction)

##      Up
## Down  0
## Up    1
```

```

glm.pred2=rep("Down", 1089)
glm.probs2=predict(glm.fit2,type="response")
glm.probs2[1:10]

##      1       2       3       4       5       6       7
## 0.5777243 0.5661029 0.5492840 0.5132426 0.6071571 0.5644982 0.5716776
##      8       9      10
## 0.5321015 0.5659641 0.5541137

glm.pred2[glm.probs2>0.5] <- "Up"
table(glm.pred2,Direction)

##          Direction
##  glm.pred2 Down Up
##      Down    33   26
##      Up      451  579

579/(579+451)

## [1] 0.5621359

451/(579+451)

## [1] 0.4378641

26/(26+33)

## [1] 0.440678

```

The confusion matrix shows that on days when logistic regression predicts an increase in the market, it has a 56.2% accuracy rate. The error rate is 43.8% for predicting Up and is actually Down. The error rate is 44.1% for predicting Down and is actually Up.

ANSWER FOR 11:

```

(a) #install.packages("RCurl")
library(RCurl)

## Loading required package: bitops

dataset<-getURL(
  'https://raw.githubusercontent.com/Jwcrist/BusA/master/homeworks/Assignment%204/Auto.csv',
  ssl.verifypeer=0L, followlocation=1L)
dataset1<-read.csv(text=dataset)
head(dataset1,3)

##   mpg cylinders displacement horsepower weight acceleration year origin
## 1 18         8           307        130   3504         12.0     70      1
## 2 15         8           350        165   3693         11.5     70      1
## 3 18         8           318        150   3436         11.0     70      1
##                           name mpg01

```

```

## 1 chevrolet chevelle malibu      0
## 2          buck skylark 320      0
## 3      plymouth satellite      0

```

As can be shown above we have created our binary data and a median column

```

(b) library(dplyr)
## Warning: package 'dplyr' was built under R version 3.1.2
##
## Attaching package: 'dplyr'
##
## The following object is masked from 'package:MASS':
##
##     select
##
## The following object is masked from 'package:stats':
##
##     filter
##
## The following objects are masked from 'package:base':
##
##     intersect, setdiff, setequal, union

dataset2<-dataset1 %>%
  select(mpg, cylinders, displacement, as.numeric(horsepower), weight, acceleration, year, origin, m

str(dataset2)

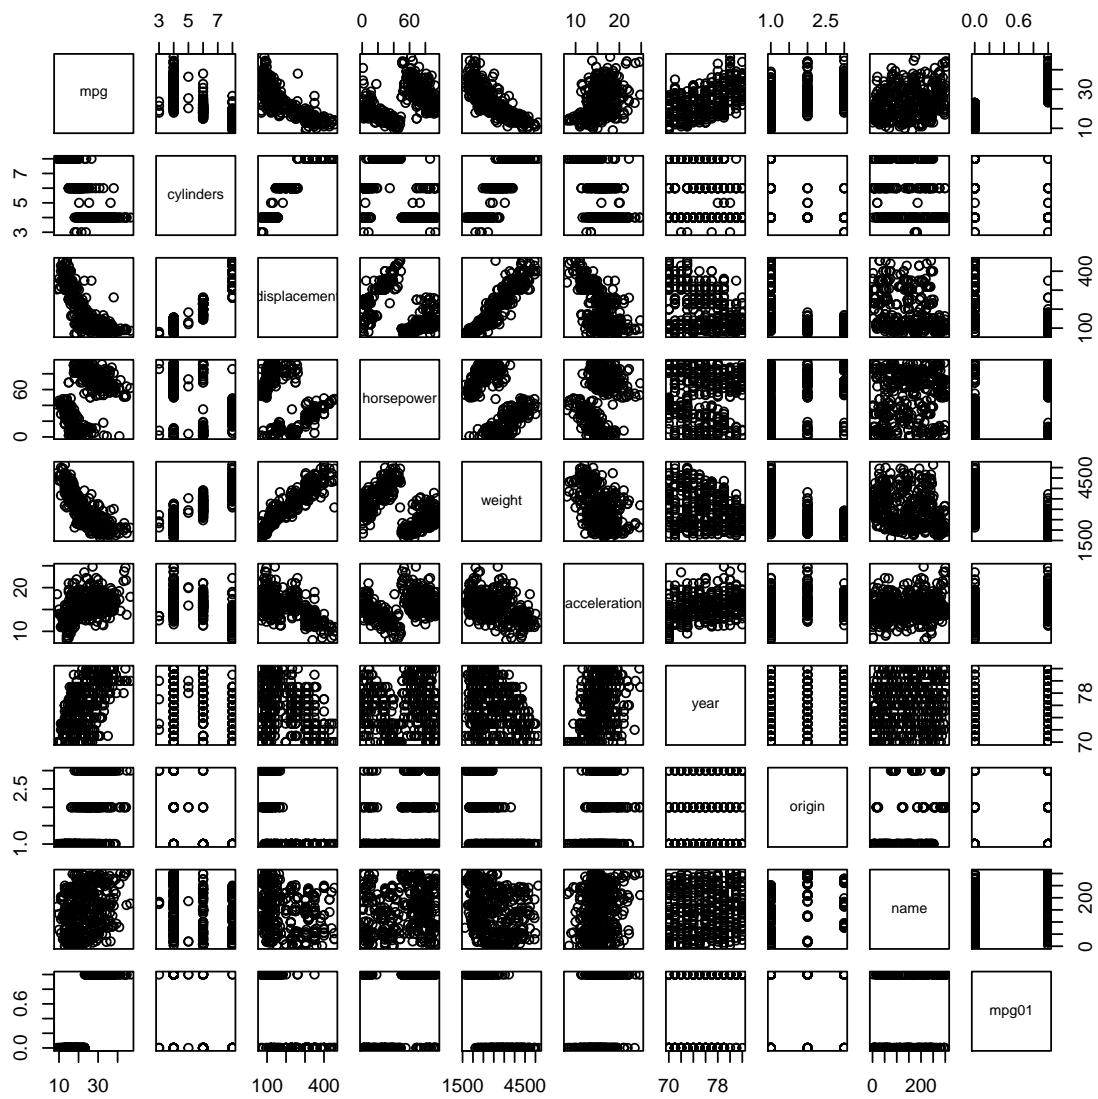
## 'data.frame': 397 obs. of  9 variables:
## $ mpg           : num  18 15 18 16 17 15 14 14 14 15 ...
## $ cylinders    : int  8 8 8 8 8 8 8 8 8 8 ...
## $ displacement: num  307 350 318 304 302 429 454 440 455 390 ...
## $ horsepower   : Factor w/ 94 levels "?","100","102",...: 17 35 29 29 24 42 47 46 48 40 ...
## $ weight        : int  3504 3693 3436 3433 3449 4341 4354 4312 4425 3850 ...
## $ acceleration: num  12 11.5 11 12 10.5 10 9 8.5 10 8.5 ...
## $ year         : int  70 70 70 70 70 70 70 70 70 70 ...
## $ origin        : int  1 1 1 1 1 1 1 1 1 1 ...
## $ mpg01        : int  0 0 0 0 0 0 0 0 0 0 ...

cor(dataset2)

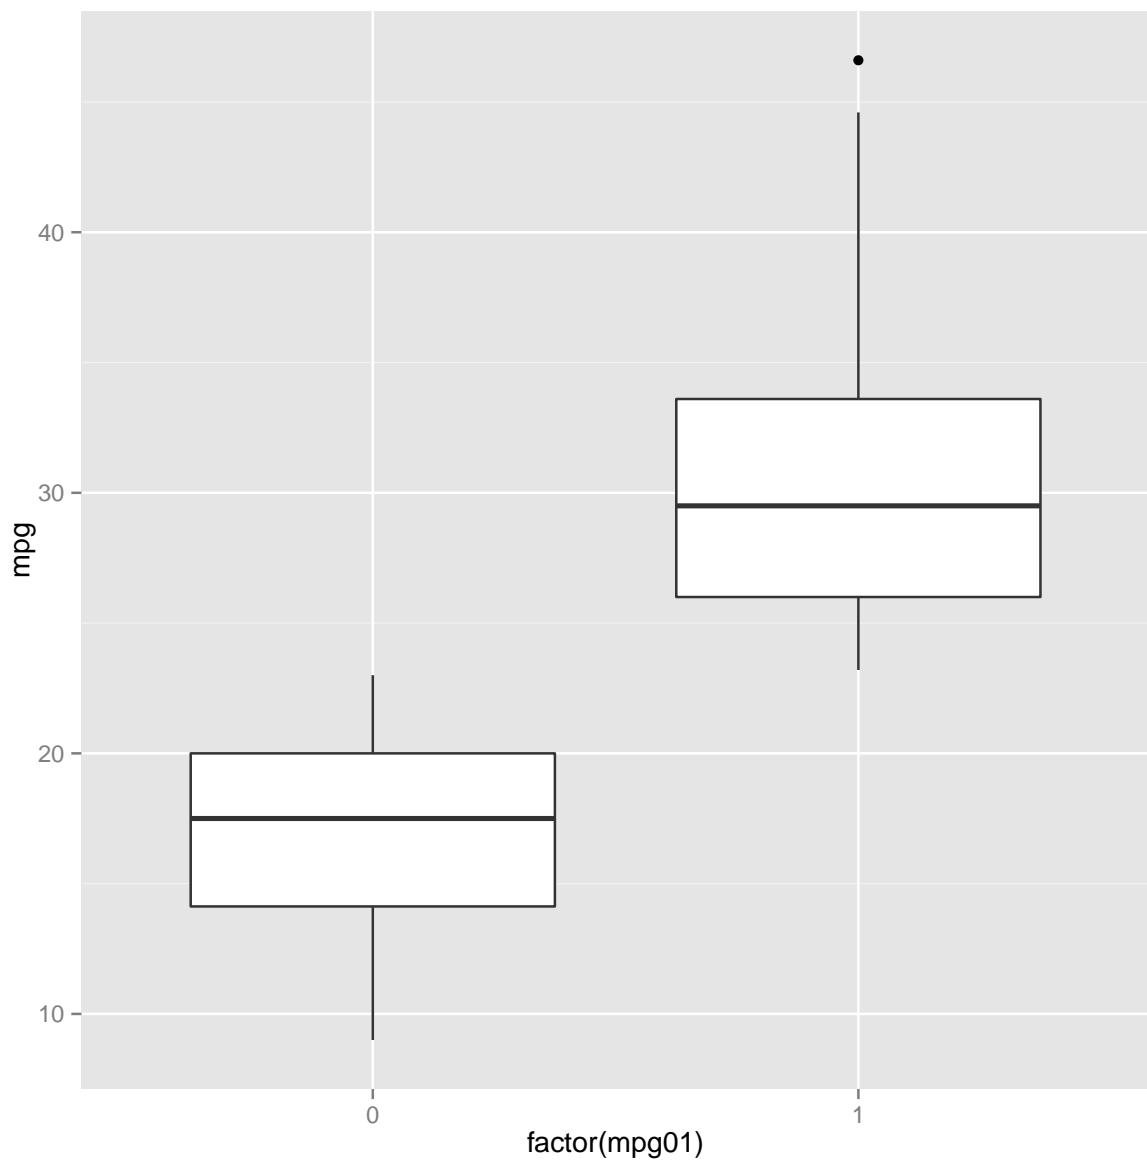
## Error in cor(dataset2): 'x' must be numeric

pairs(dataset1)

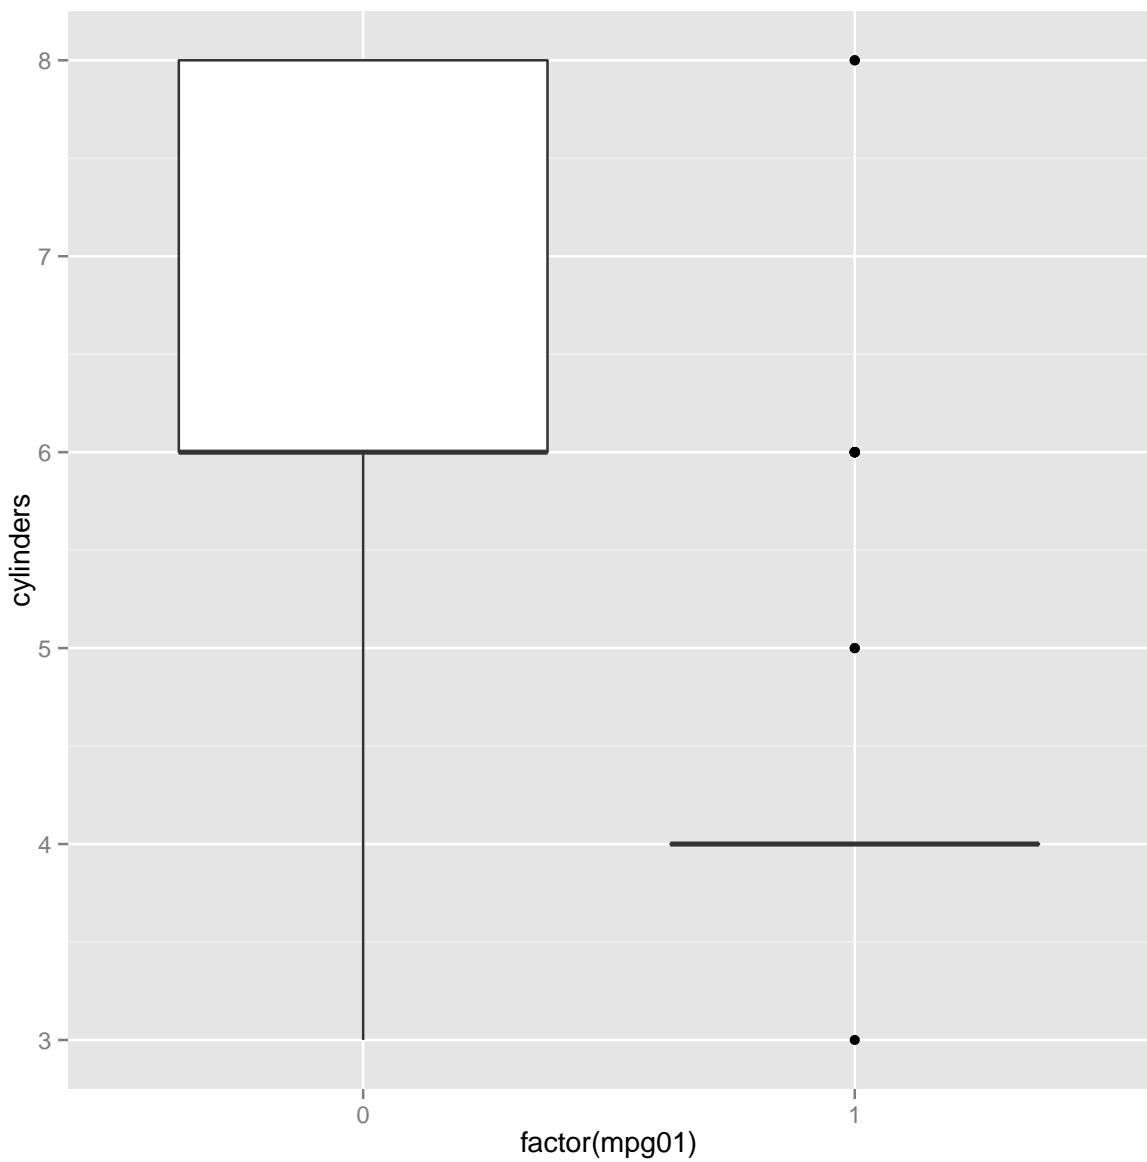
```



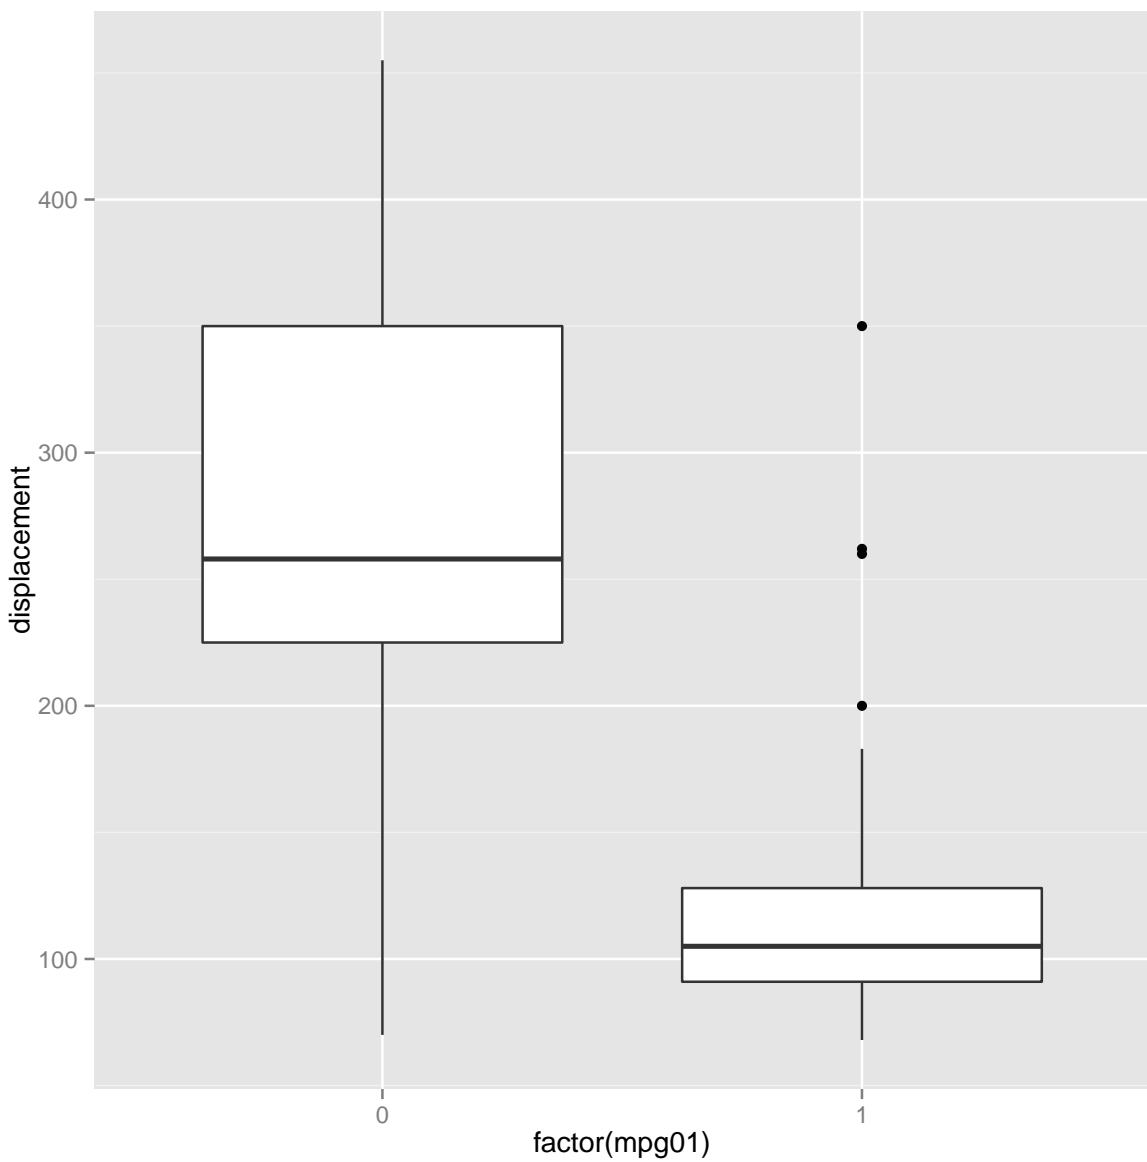
```
library(ggplot2)
ggplot(dataset1, aes(x=factor(mpg01), y=mpg))+geom_boxplot()
```



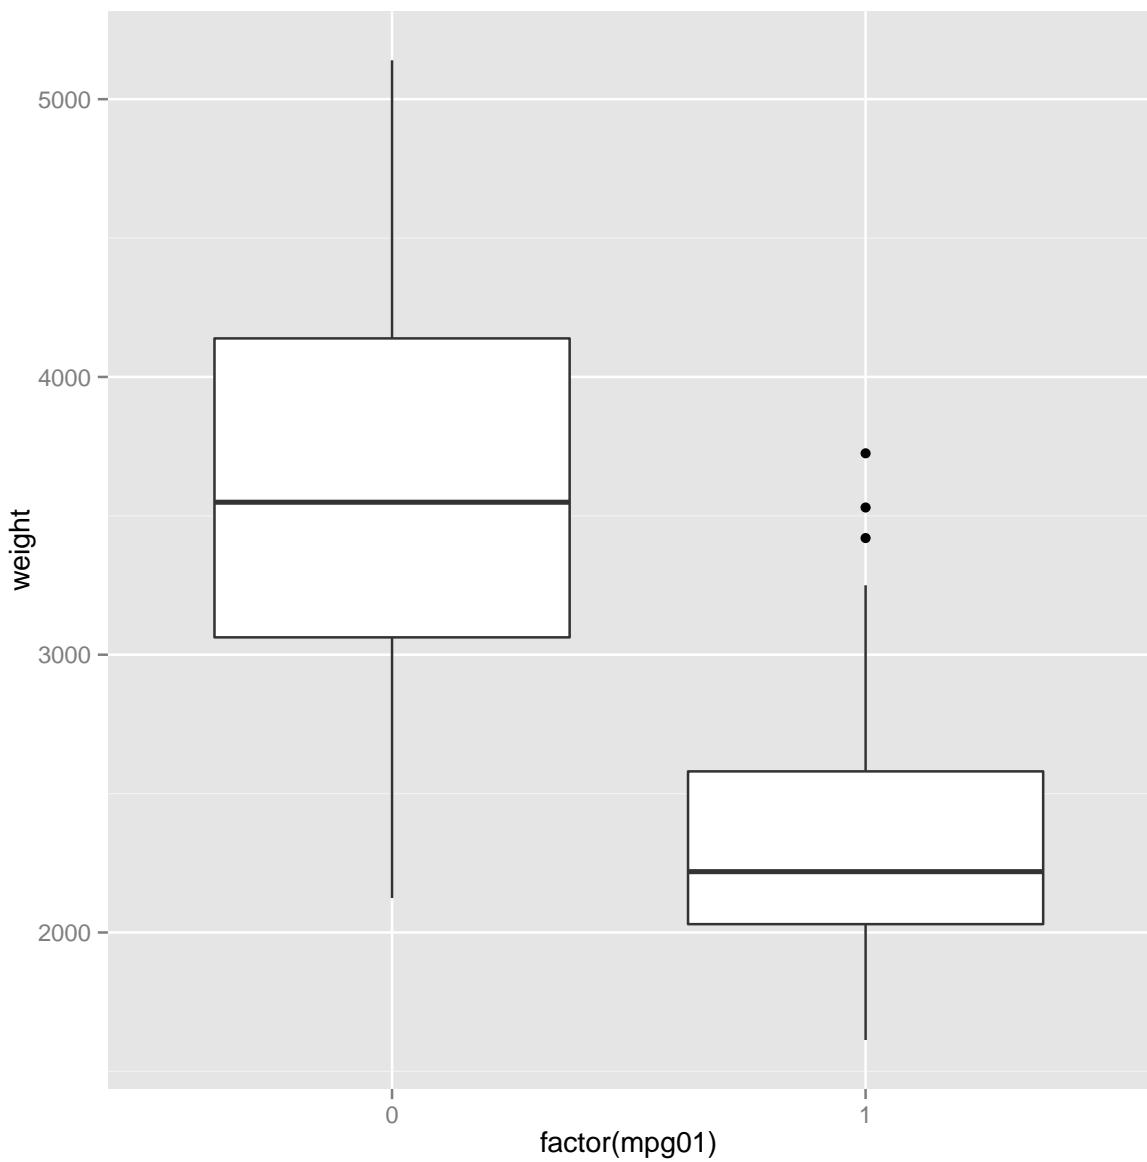
```
ggplot(dataset1, aes(x=factor(mpg01), y=cylinders)) + geom_boxplot()
```



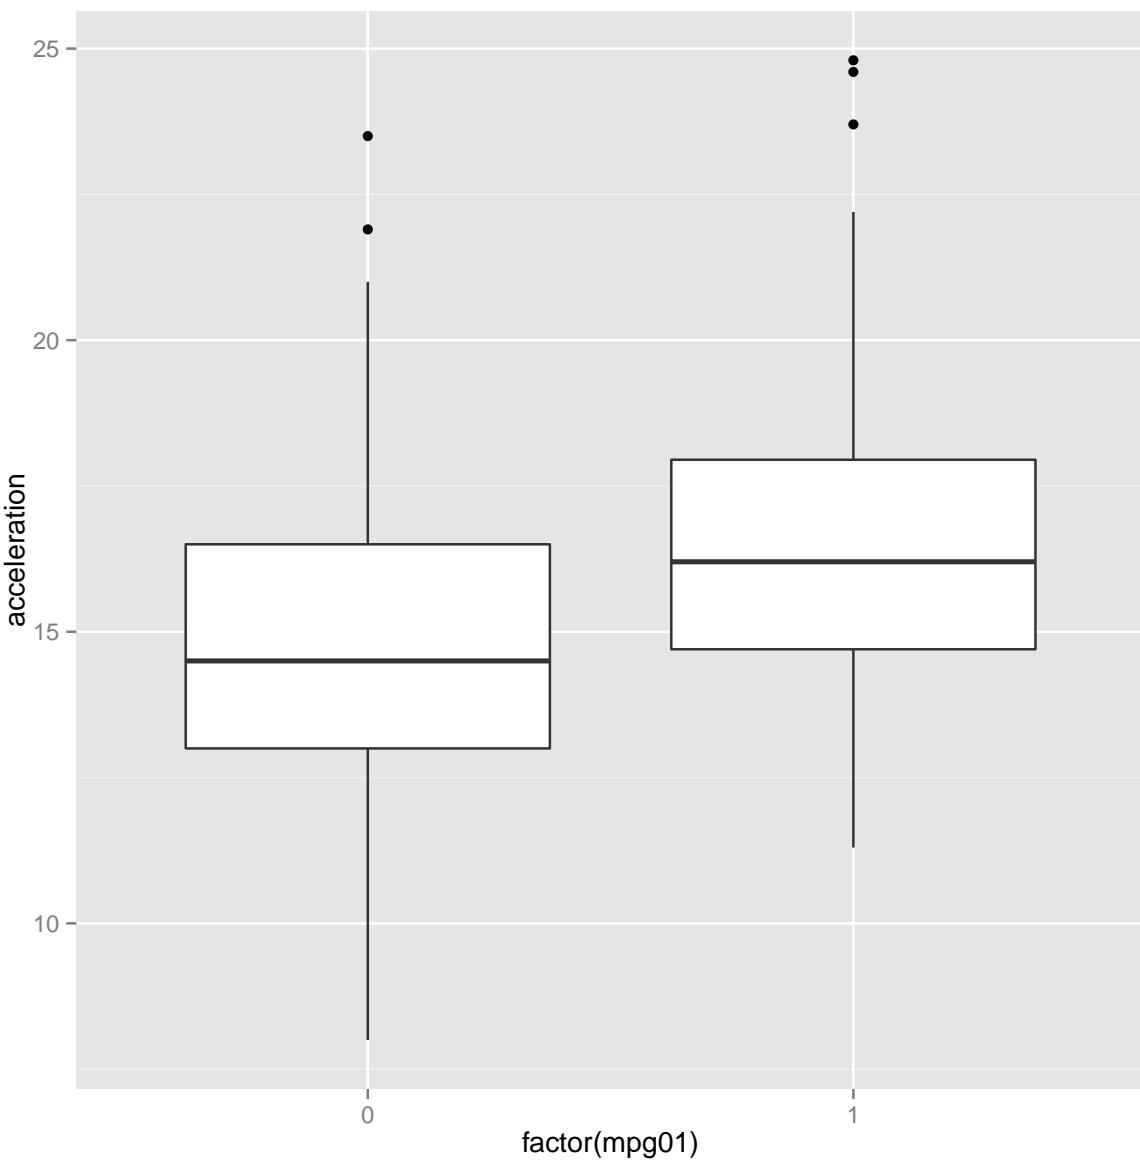
```
ggplot(dataset1, aes(x=factor(mpg01), y=displacement))+geom_boxplot()
```



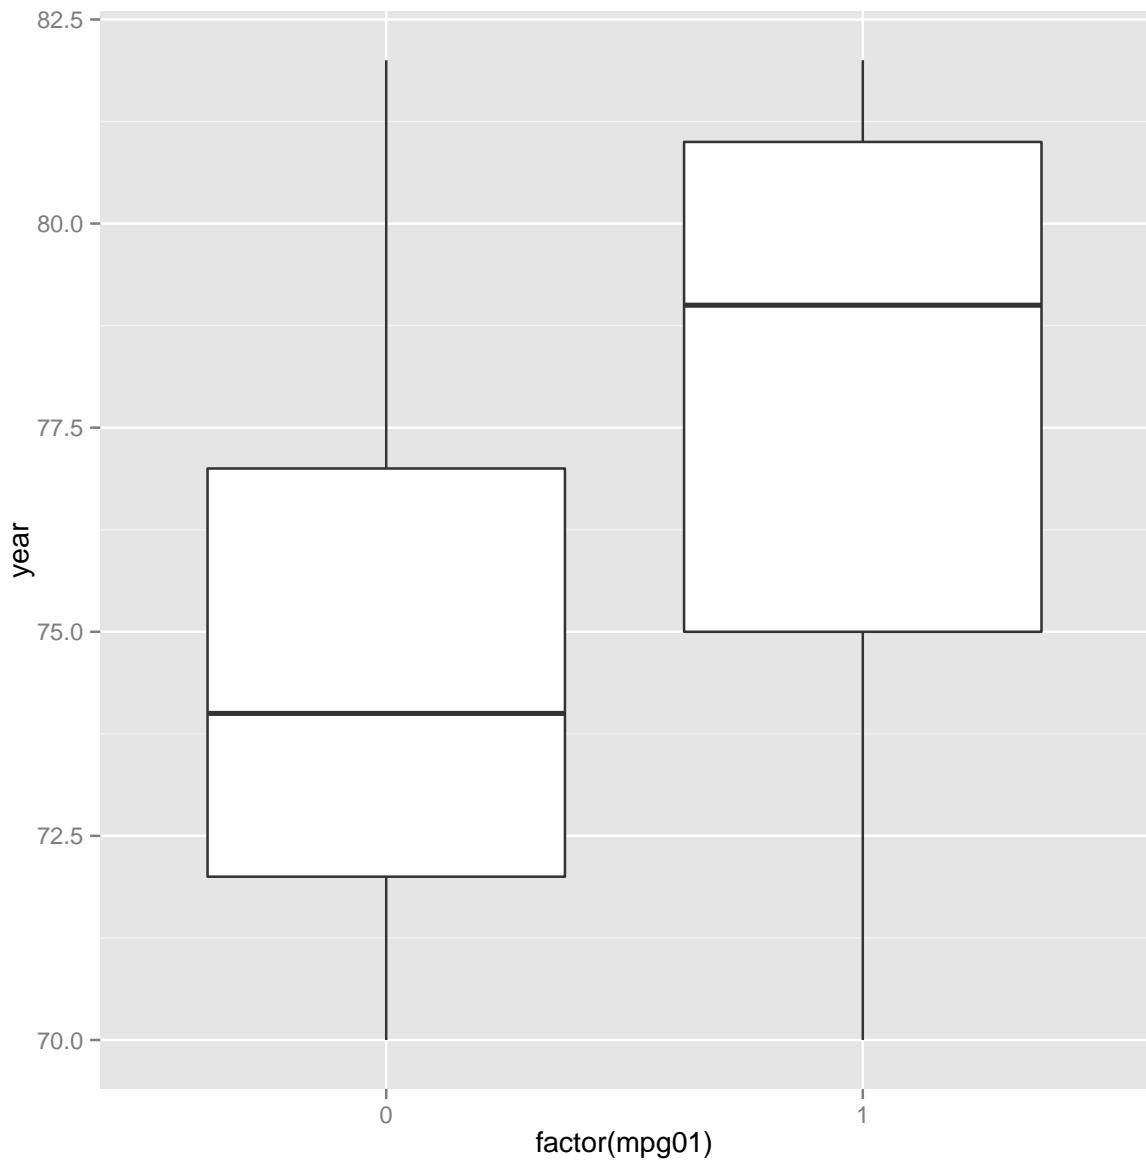
```
ggplot(dataset1, aes(x=factor(mpg01), y=weight)) + geom_boxplot()
```



```
ggplot(dataset1, aes(x=factor(mpg01), y=acceleration))+geom_boxplot()
```



```
ggplot(dataset1, aes(x=factor(mpg01), y=year))+geom_boxplot()
```



There is too many layers to observe anything of significance in horsepower. Acceleration and year may have be significantly different from eachother. All the other plots show that there is probable differences between our binary factors.

(c) `tail(dataset1)`

```
##      mpg cylinders displacement horsepower weight acceleration year origin
## 392    27          4           151         90    2950        17.3     82      1
## 393    27          4           140         86    2790        15.6     82      1
## 394    44          4           97          52    2130        24.6     82      2
## 395    32          4           135         84    2295        11.6     82      1
## 396    28          4           120         79    2625        18.6     82      1
## 397    31          4           119         82    2720        19.4     82      1
##                               name mpg01
## 392  chevrolet camaro      1
```

```

## 393  ford mustang gl      1
## 394      vw pickup      1
## 395      dodge rampage    1
## 396      ford ranger      1
## 397      chevy s-10       1

trainingdata<-dataset1[1:318,]
testdata<-dataset1[79,]

```

A 80% training data set was chosen to be used in order to prove against our test data.

(d) Do Not Do

(e) Do Not Do

```

(f) glm.fit3=glm(mpg01~Lag2,data=Weekly, family = binomial)
## Error in eval(expr, envir, enclos): object 'mpg01' not found

summary(glm.fit2)

##
## Call:
## glm(formula = Direction ~ Lag2, family = binomial, data = Weekly)
##
## Deviance Residuals:
##     Min      1Q  Median      3Q      Max
## -1.564  -1.267   1.008   1.086   1.386
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.21473   0.06123   3.507 0.000453 ***
## Lag2         0.06279   0.02636   2.382 0.017230 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 1496.2 on 1088 degrees of freedom
## Residual deviance: 1490.4 on 1087 degrees of freedom
## AIC: 1494.4
##
## Number of Fisher Scoring iterations: 4

coef(glm.fit2)

## (Intercept)      Lag2
## 0.21473151  0.06279058

contrasts(Direction)

##      Up
## Down  0
## Up    1

```

```

glm.pred2=rep("Down", 1089)
glm.probs2=predict(glm.fit2,type="response")
glm.probs2[1:10]

##      1       2       3       4       5       6       7
## 0.5777243 0.5661029 0.5492840 0.5132426 0.6071571 0.5644982 0.5716776
##      8       9      10
## 0.5321015 0.5659641 0.5541137

glm.pred2[glm.probs2>0.5] <- "Up"
table(glm.pred2,Direction)

##          Direction
##  glm.pred2 Down  Up
##      Down    33   26
##      Up     451  579

```